

R 101, v2

Nastasiya F. Grinberg

03-12-2025

Here, we review some fundamental principles coding in R, such as different object types, subsetting and indexing, writing functions, etc. These notes are by no means exhaustive, but they set out the basics that will help you start with R-based data analysis and manipulation.

Vectors, Matrices, Lists and Data Frames

We start by exploring the ways you can store and manipulate data in R.

Vectors

A vector is a collection of elements of a given type (e.g. numeric, or character). A vector has only one dimension—its length. You can check a vector's length using the `length` function. There are several ways to create vectors in R:

- Using the `c` function, like so `c(1, 10)`.
- Consecutive numbers using the `:` operator, e.g. `1 : 10, (-10) : 2`.
- Using the functions `rep` and `seq`:
 - `rep(NA, 10)` creates a vector of NAs of length 10.
 - `seq(0, 1, by = 0.1)` creates a vector of values from 0 to 1 with step 0.1.
 - `seq(0, 1, len = 20)` creates a vector of evenly spaced values from 0 to 1 of length 20.
- You can combine all of the above, e.g. `c(1 : 4, seq(0, 1, by = 0.1), 2)`.

You can extract elements from vectors using the square brackets `[]`, e.g. `vec[c(1, 3)]` would return the 1st and 3rd element of `vec`; `vec[1:10]` would return the first 10 elements of `vec`. To *omit* elements of a vector, use *negative indices*, e.g. `vec[-1]` returns all *but the 1st element* of `vec`. Note that indexing in R *starts from 1*, not 0 as in many other programming languages. This is true not only for vectors but for all other types of objects that have dimensions.

Combining conditions and the square brackets is an effective way to extract vector elements with particular properties, e.g. `vec[vec > 10]` returns all elements of `vec` greater than 10, and `vec[vec == 0 | vec > 5]` would return all elements that are either equal to 0 or are bigger than 5.

Matrices

Matrices are rectangular arrays holding elements of a given type. As such, they have two dimension parameters: number of rows and number of columns. You can check matrix dimensions using the `nrow`, `ncol`, and `dim` functions. Matrices are usually created using the `matrix` function:

- `matrix(1 : 10, nrow = 2)` creates a matrix from values `1 : 10` arranged in a 2×5 grid. By default, values would be arranged by column. If you want to change this behaviour, use the `byrow = TRUE` argument.
- `matrix(NA, nrow = 2, ncol = 5)` creates a matrix of `NAs` with 2 rows and 5 columns.

Extracting matrix elements is also performed with `[]` and using the same principles as for vectors. The only difference being that, since matrices are two-dimensional objects, you need to provide indices for both dimensions, separated by a comma, within the square brackets: `[row_indices, col_indices]`. For example, `mat[1:10, 1:10]` returns the submatrix consisting of the first ten rows and first ten columns of the original matrix. Leaving one dimension empty returns all entries of that dimension, e.g. `mat[1 : 10,]` would return the submatrix of the first ten rows and *all* the columns.

Lists

A list is another one-dimensional structure but, unlike vectors, they can hold data of different types and are thus very useful. You can create a named list like so:

```
my_list <- list(a = 10, b = 'foo', c = list(1, 2, 3), d = matrix(1 : 20, nrow = 4))
my_list
```

```
## $a
## [1] 10
##
## $b
## [1] "foo"
##
## $c
## $c[[1]]
## [1] 1
##
## $c[[2]]
## [1] 2
##
## $c[[3]]
## [1] 3
##
##
## $d
##      [,1] [,2] [,3] [,4] [,5]
## [1,]     1     5     9    13    17
## [2,]     2     6    10    14    18
## [3,]     3     7    11    15    19
## [4,]     4     8    12    16    20
```

Note that providing names is not necessary (try `list(10, 'foo')`). Indexing lists can be done with double square brackets and a numerical index, with the `$` operator, or with double square brackets and an element name:

```
my_list[[2]] # second entry
```

```
## [1] "foo"
```

```

my_list$a # an entry called "a"

## [1] 10

my_list[["b"]] # an entry called "b"

## [1] "foo"

```

If you have a list of single values of the same class, sometimes you might want to “unlist” it into a vector:

```

my_list <- list(1, 2, 3)
my_list

```

```

## [[1]]
## [1] 1
##
## [[2]]
## [1] 2
##
## [[3]]
## [1] 3

```

```
unlist(my_list)
```

```
## [1] 1 2 3
```

Data Frames

Like matrices, data frames are rectangular arrays of data but, unlike matrices, data frames can hold values of more than just one type. Data frames (rather than matrices) are customarily used to store data for data analysis. You can create the `data.frame` from scratch:

```

my_df <- data.frame(col1 = 1 : 10, col2 = 11 : 20)
my_df

```

```

##      col1 col2
## 1      1   11
## 2      2   12
## 3      3   13
## 4      4   14
## 5      5   15
## 6      6   16
## 7      7   17
## 8      8   18
## 9      9   19
## 10     10  20

```

or convert a matrix:

```

mat <- matrix(1 : 10, nrow = 2)
mat

##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    3    5    7    9
## [2,]    2    4    6    8   10

as.data.frame(mat)

##   V1 V2 V3 V4 V5
## 1  1  3  5  7  9
## 2  2  4  6  8 10

data.frame(mat)

##   X1 X2 X3 X4 X5
## 1  1  3  5  7  9
## 2  2  4  6  8 10

```

However, more likely than not you would be exporting your data into R using a `read.csv` function.

For the following, to illustrate indexing and subsetting, we are using the inbuilt dataset `mtcars`. You can extract a subsets of a `data.frame` using using the same bracket syntax as for matrices, e.g.:

```

mtcars[1, 1] # returns the first entry of the first column

## [1] 21

mtcars[1 : 4, 5] # returns the first four entries (rows) for the fifth column

## [1] 3.90 3.90 3.85 3.08

```

Additionally, you can also extract columns using their names (lists style):

```

mtcars$carb # extract the column called "carb"

## [1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6 8 2

mtcars[["carb"]] # extract the column called "carb"

## [1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6 8 2

```

You can also combine approaches, e.g.:

```

mtcars[c(1, 2, 5), "carb"] # extract entries (rows) 1, 2, and 5 of column "carb"

## [1] 4 4 2

```

Vectorisation

Many operations on vectors and matrices that would require explicit loops in other languages are vectorised in R and can be applied directly to vector and matrix objects. For example:

- For a matrix `mat`, `mat^2` returns a matrix with every entry of the original matrix squared. Similarly, `sqrt(mat)` returns a matrix with entries being square roots of the corresponding entries of `mat`.
- For two matrices of the same dimensions `mat1` and `mat2`, `mat1 + mat2` creates a matrix with each entry being a sum of the corresponding entries of the first two matrices.
- For a vector `vec`, `vec^2` returns a vector with every entry of the original vector squared. Similarly, `sqrt(vec)` returns a vector with entries being square roots of entries of `vec`.
- For two vectors of the same length `vec1` and `vec2`, `vec1 + vec2` creates a vector with each entry being a sum of the corresponding entries of the first two vectors. Note that if vectors are of different lengths, the shorter vector is *recycled* (try it!).
- Both for vectors and for matrices, conditions, e.g. `mat == 0` or `vec == 0`, returns a vector/matrix of the same dimensions with `TRUE` for every entry satisfying a condition, and `FALSE` otherwise. Try, e.g. `mat <- matrix(1 : 10, nrow = 2)` and `mat > 3` or `vec <- rnorm(10)` and `vec > 0`.
- Note that `TRUE` and `FALSE` correspond to numerical values of 1 and 0, respectively. Thus, the `sum` function can be effectively combined with conditions for counting, e.g. `sum(vec > 0)` counts a number of positive elements of `vec`.

Common and useful built-in functions

- `colSums` and `rowSums` for summing columns and rows of a matrix or a data frame. Use the argument `na.rm = TRUE` if your matrix contains `NA` in order to ignore missing values.
- `colMeans` and `rowMeans` for finding means of columns and rows of a matrix or a data frame. Use the argument `na.rm = TRUE` if your matrix contains `NA` in order to ignore missing values.
- `sample` can be used to sample from a vector of values with or without replacement, e.g. `sample(1 : 10, 100, replace = TRUE)` returns 100 values randomly chosen from 1:10 with replacement (i.e. any number can be chosen more than once) or `sample(1 : 5)` returns 5 values from 1:5 without replacement (i.e. a random permutation of 1:5).
- `is.na(x)` returns `TRUE` if `x` is `NA` and `FALSE` otherwise. `is.na` is vectorised, which means that it works on vectors and matrices by applying itself to each of its elements individually.
- `which.min` and `which.max`, when applied to a vector, return the index of its minimum and maximum element, respectively. For example, `which.max(c(2, 10, 5))` would return 2 because 10 is the largest element of the vector argument.
- `unique`, given a vector, returns a vector of unique values featuring in the input: try `unique(c(1, 2, 3, 1, 2))` and `unique(iris$Species)`.
- `table`, given a vector of values, outputs a table of counts of occurrences of each unique value in the vector: try `table(iris$Species)`. Given two or more vectors, `table` counts the number of occurrences of all unique tuples (e.g. n -way combinations, where n is the number of vectors provided), e.g. `table(mtcars$cyl, mtcars$gear)`.

Loops

Remember that if you calculate values within a loop, which you would like to save, you need to create an object to store these values. For example, if you want to create a loop to square numbers from 1 to 10 (obviously, a much easier way to do this is `(1 : 10)^2`), you need to create an empty vector to store the squares first:

```

res <- rep(NA, 10)
for(i in 1 : 10) {
  res[i] <- i^2
}
print(res)

## [1] 1 4 9 16 25 36 49 64 81 100

```

Similarly, if you would like to create a double loop to calculate i^k for $i = 1, \dots, 5$ and $k = 1, 2, 3$, you would need a 5×3 (or a 3×5) matrix for your results:

```

res_mat <- matrix(NA, nrow = 5, ncol = 3)
for(i in 1 : 5) {
  for(k in 1 : 3) {
    res_mat[i, k] <- i^k
  }
}
print(res_mat)

```

```

##      [,1] [,2] [,3]
## [1,]     1     1     1
## [2,]     2     4     8
## [3,]     3     9    27
## [4,]     4    16    64
## [5,]     5    25   125

```

*apply functions

The family of ***apply** functions (most common of which are **apply**, **sapply**, and **lapply**) are an alternative to loops and a very R-esque way of applying a function to elements of a matrix, vector, or list. The functions you can use within the ***apply** functions, just as within loops, can be arbitrarily complex.

- **apply** can be used to apply a function to each row (1) or column (2) of a matrix. For example, `apply(mat, 1, max)` will find the biggest element in each row (argument 1) and return them as a vector. You can also supply a more sophisticated function to be applied, e.g. `apply(mat, 2, FUN = function(x) {x[-which.min(x)]})` returns each column *except* the minimum value.
- **lapply** and **sapply** apply a function to each element of a list/vector and return the output in the form of a list or a vector, respectively. Try `lapply(1 : 10, FUN = function(x) {x^2})` and `sapply(1 : 10, FUN = function(x) {x^2})`. The above double-loop can be rewritten using the **sapply** function like so:

```

sapply(1 : 5, FUN = function(i) {
  sapply(1 : 3, FUN = function(k) {
    i^k
  })
})

```

```

##      [,1] [,2] [,3] [,4] [,5]
## [1,]     1     2     3     4     5
## [2,]     1     4     9    16    25
## [3,]     1     8    27    64   125

```

The advantage of `sapply/lapply` over loops is that you do not need to create a dedicated container for the output.

Writing your own functions

Most likely, built-in functions would not be enough to cater for all your needs and more often than not you will have to write your own bespoke functions. To create a function, you follow the following template:

```
fun <- function(x, y) {  
  # body of the function  
  # return(output)  
}
```

where `fun` is the name of your function and `x` and `y` are function arguments. Note, however, that `return` is optional because, by default, your function returns the value of the last expression it evaluates. So `fun1 <- function(x) {x^2}` and `fun2 <- function(x) {ans <- x^2; return(ans)}` will return the same answer. Your function can have as many arguments as you like; they might or might not have default values.

Here is an example of a function that calculates sample variance for a vector of values `x` (of course there is an in-built function that does that called `var`):

```
my_var <- function(x) {  
  xbar <- mean(x)  
  sum((x - xbar)^2)/(length(x) - 1)  
}  
  
my_var(rnorm(10))  
  
## [1] 0.6446833
```

Here is an example of a function with two arguments, one of which has a default value:

```
sum_of_powers <- function(x, k = 2) {  
  sum(x^k)  
}  
  
sum_of_powers(rep(1, 10)) # uses default value of k = 2  
  
## [1] 10  
  
sum_of_powers(1 : 4, 3) # k is set to 3  
  
## [1] 100
```

Useful Links

- General introduction to R: “An Introduction to R” (Douglas, Mancini, Couto & Lusseau)
- Introduction to R with `tidyverse` and with application to data analysis: “R for Data Science” (Wickham & Grolemund)
- An Introduction to Statistical Learning with Applications in R: ITSL (James, Witten, Hastie & Tibshirani)